



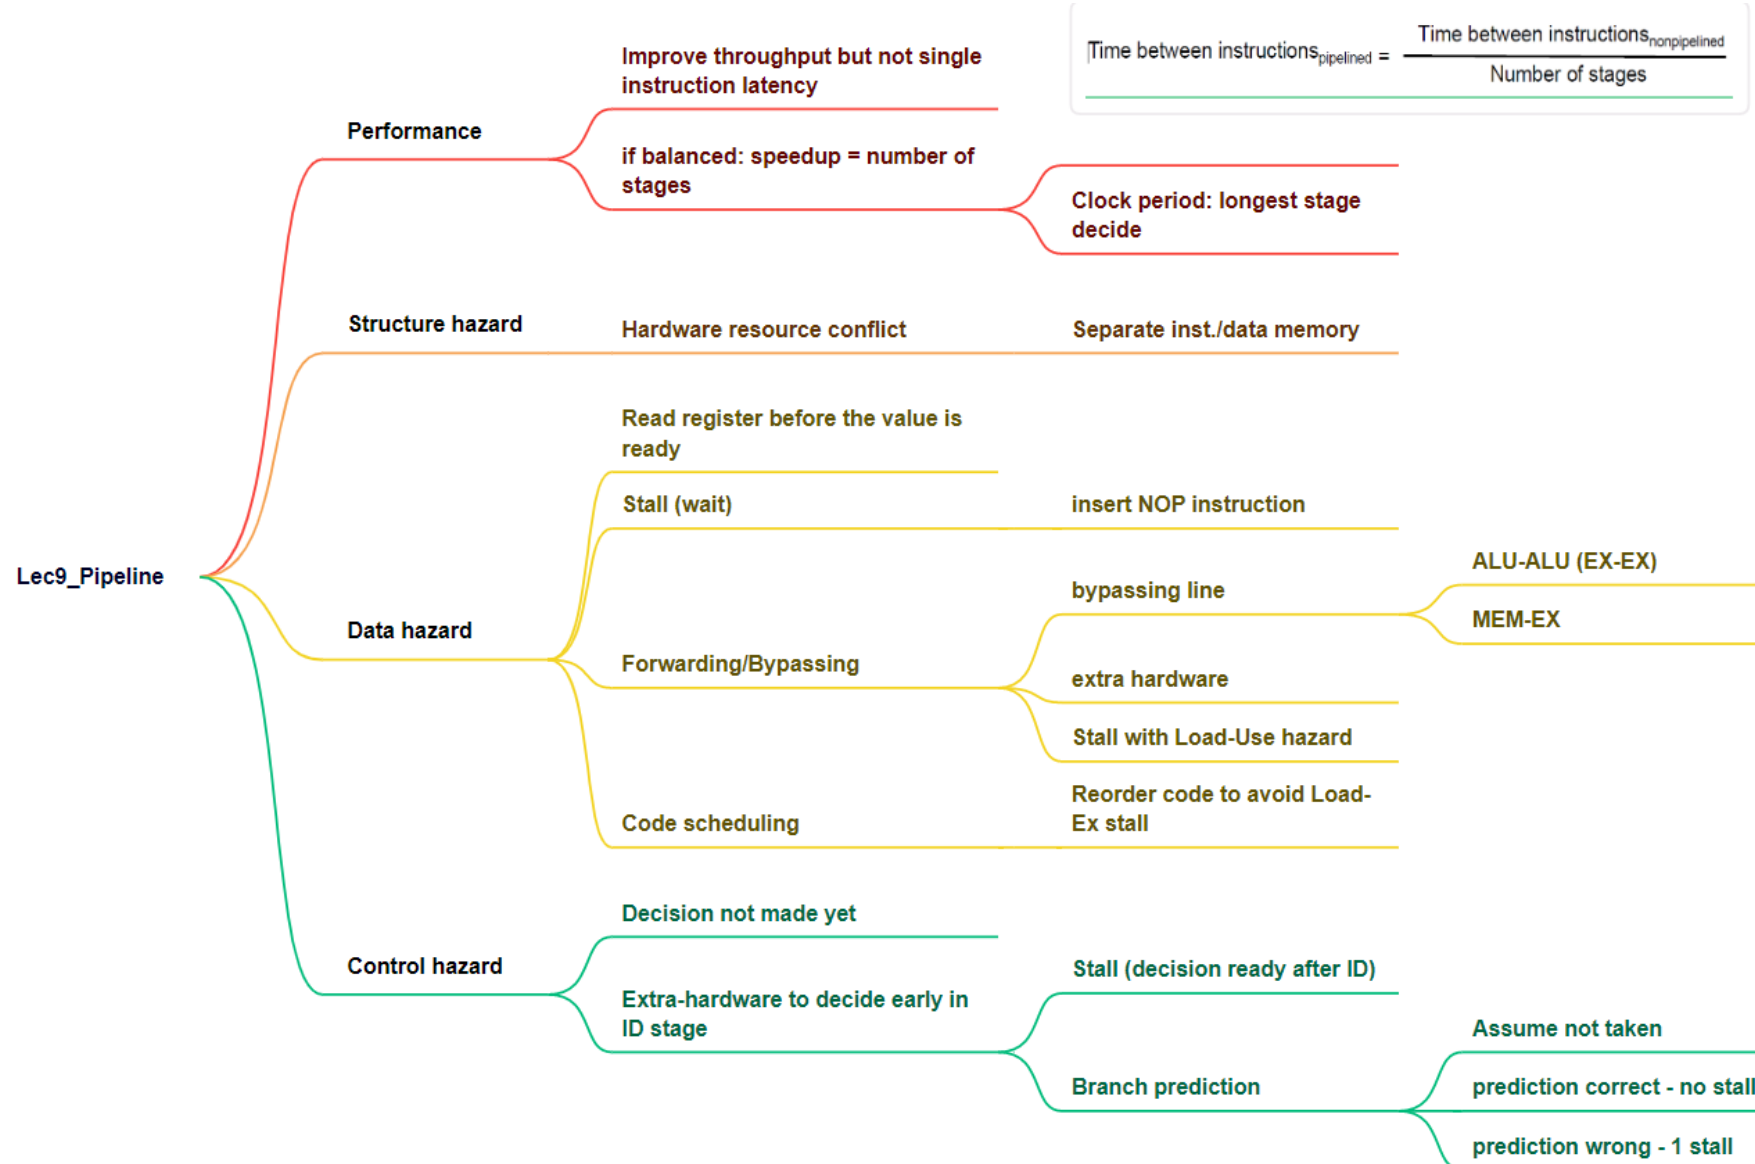
COMPUTER ORGANIZATION

Lecture 10 Instruction-Level Parallelism

2025 Spring


This PowerPoint is for internal use only at Southern University of Science and Technology. Please do not repost it on other platforms without permission from the instructor.

Recap



Instruction-Level Parallelism (ILP)

- Instruction-level parallelism: parallelism among instructions
 - Pipelining is one type of ILP: because pipeline executes multiple instructions in parallel
- To increase ILP
 - Deeper pipeline
 - Less work per stage \Rightarrow shorter clock cycle
 - Multiple issue(多发射)
 - Replicate pipeline datapath \Rightarrow multiple pipelines
 - Start multiple instructions per clock cycle
 - $CPI_{\text{overall}} < 1$, so use Instructions Per Cycle (IPC)
 - E.g., 4-way multiple-issue
 - peak $CPI = 0.25$, peak $IPC = 4$
 - But dependencies reduce this in practice



Multiple Issue

- Static multiple issue (decision made by **compiler**)
 - Compiler groups instructions to be issued together
 - Packages them into “issue slots”
 - Compiler detects and avoids hazards
- Dynamic multiple issue (decision made by **processor**)
 - CPU examines instruction stream and chooses instructions to issue each cycle
 - Compiler can help by reordering instructions
 - CPU resolves hazards using advanced techniques at runtime

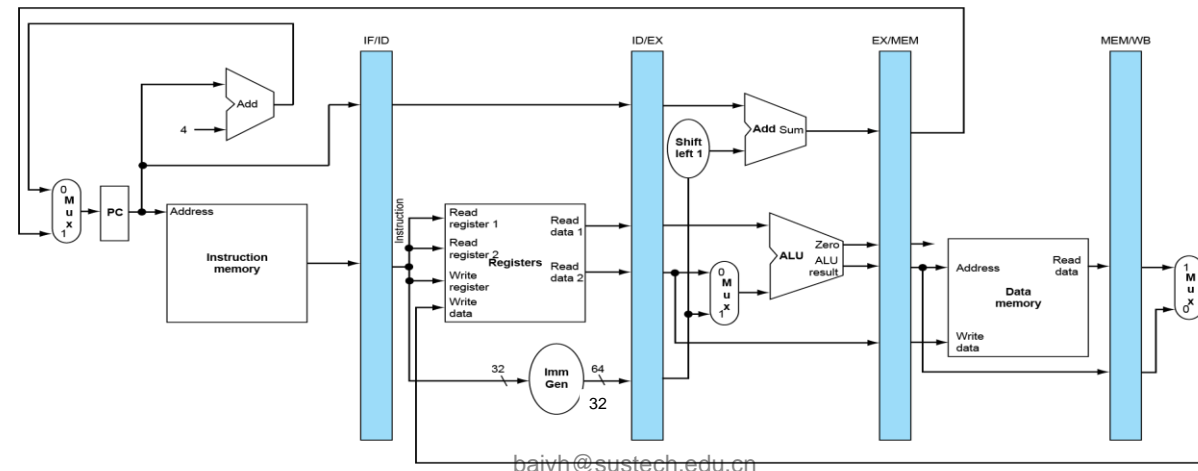
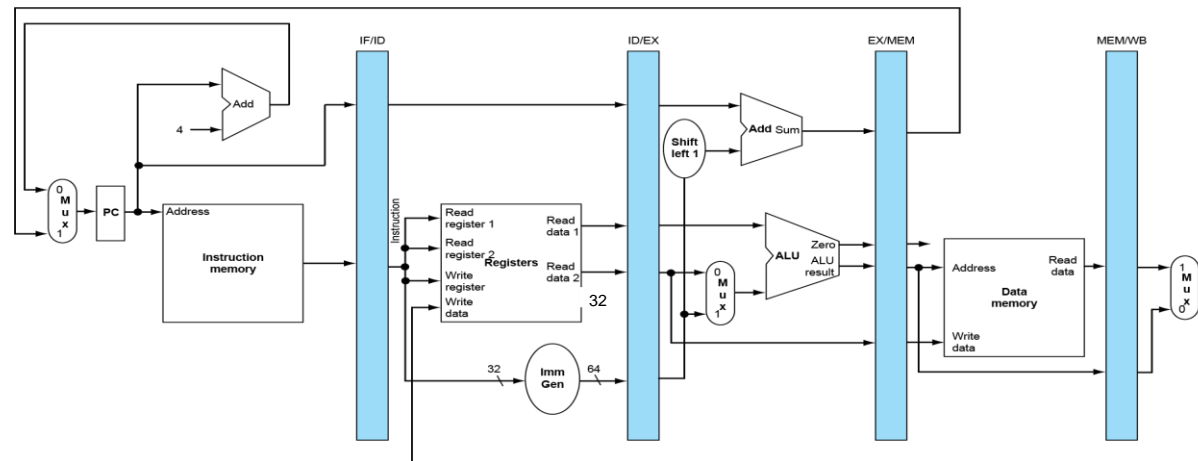


Static Multiple Issue

- Compiler groups instructions into “**issue packets**”
 - Group of instructions that can be issued on a single cycle
 - Determined by pipeline resources required
- Think of an issue packet as a very long instruction
 - Specifies multiple concurrent operations
 - \Rightarrow Very Long Instruction Word (VLIW)

Naïve Static Dual Issue

- A naïve solution: replicate the datapath to issue 2 instructions per cycle.

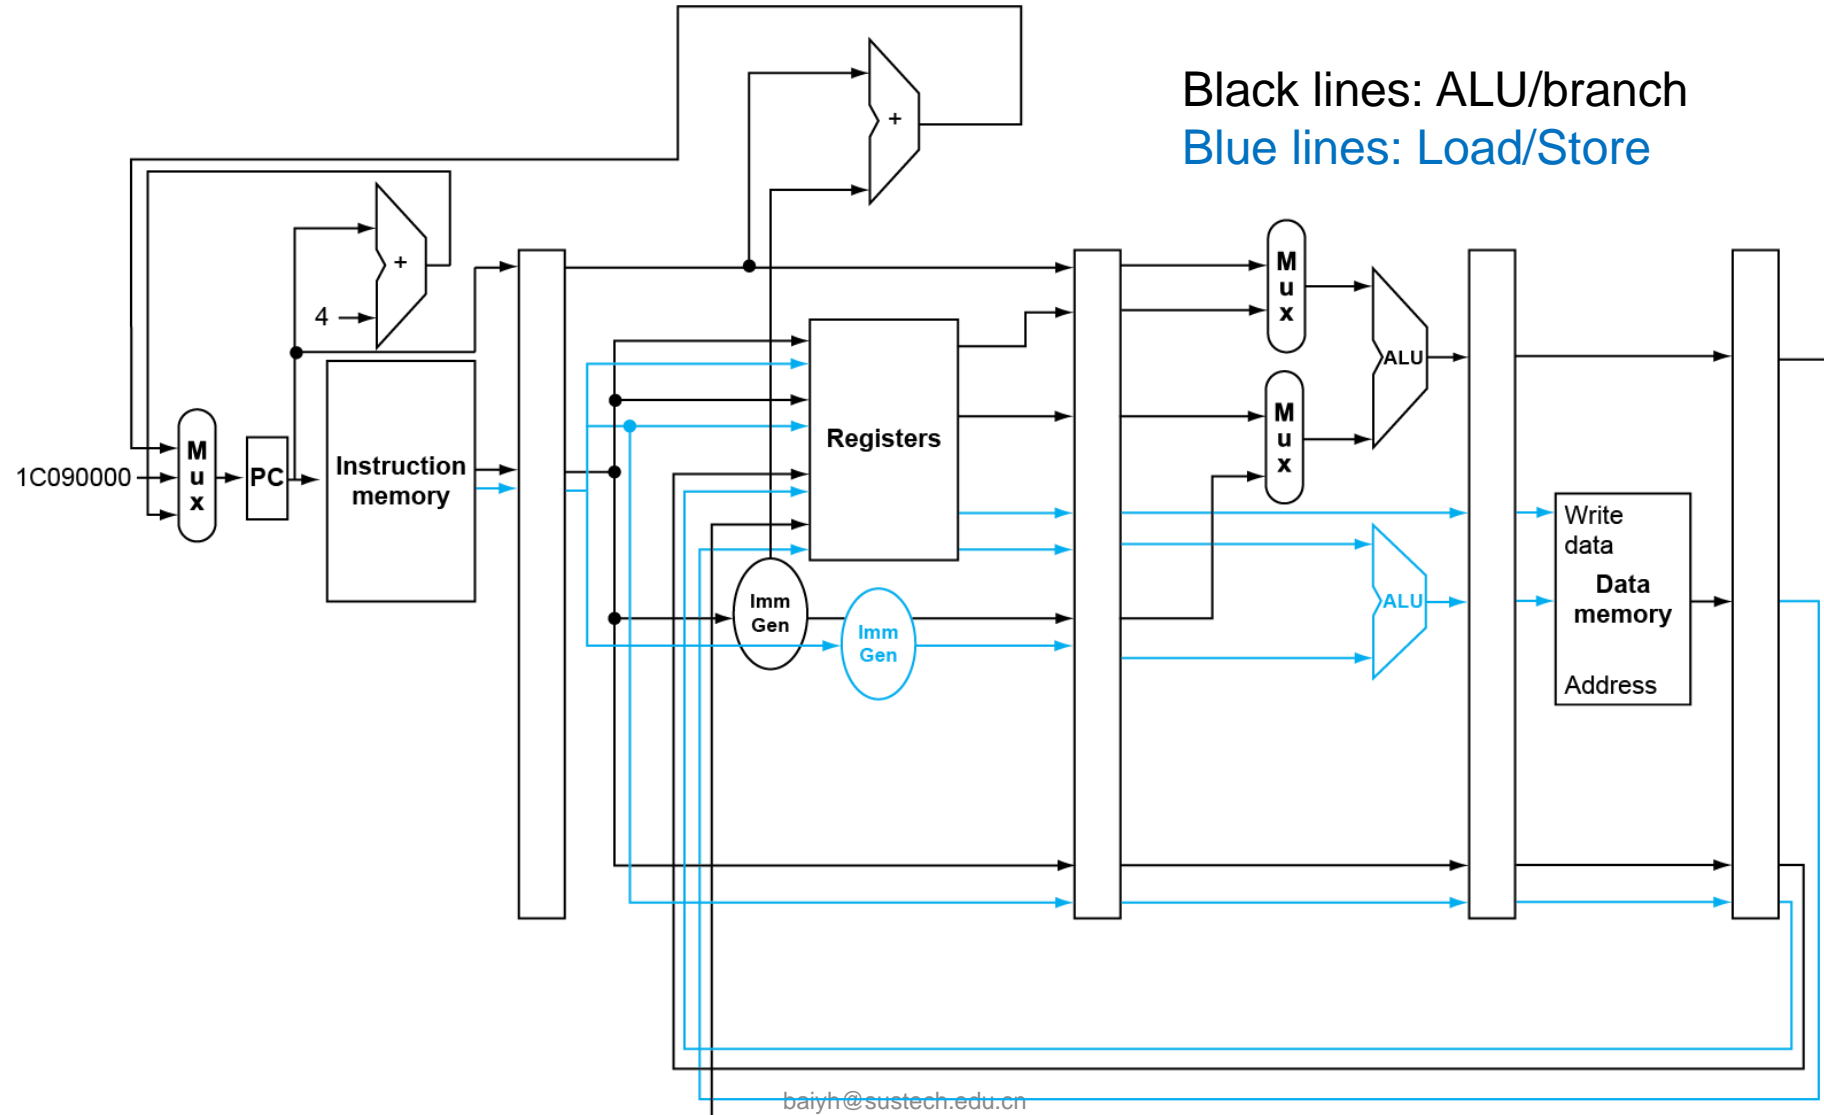


RISC-V with Static Dual Issue

- Two-issue packets
 - Divide instructions into two types:
 - Type1: ALU/branch instruction
 - Type2: load/store instruction
 - During each cycle, execute type1 and type2 inst. simultaneously
 - 64-bit aligned instructions
 - Pad an unused instruction with nop

Address	Instruction type	Pipeline Stages						
n	ALU/branch	IF	ID	EX	MEM	WB		
n + 4	Load/store	IF	ID	EX	MEM	WB		
n + 8	ALU/branch		IF	ID	EX	MEM	WB	
n + 12	Load/store		IF	ID	EX	MEM	WB	
n + 16	ALU/branch			IF	ID	EX	MEM	WB
n + 20	Load/store			IF	ID	EX	MEM	WB

RISC-V with Static Dual Issue



Hazards in the Dual-Issue RISC-V

- More instructions executing in parallel
- EX data hazard
 - Forwarding avoided stalls with single-issue
 - Now can't use ALU result in load/store in same packet
 - add `t0, s0, s1`
lw `s2, 0(t0)`
 - Split into two packets, effectively a stall
- Load-use hazard
 - Still one cycle **use latency** (number of clock cycles between load and use), but now two instructions
- More aggressive scheduling required



Scheduling Static Multiple Issue

- Compiler must remove some/all hazards
 - Reorder instructions into issue packets
 - No dependencies within a packet
 - Possibly some dependencies between packets
 - Varies between ISAs; compiler must know!
 - Pad with nop if necessary

Scheduling Example

- Schedule this for dual-issue RISC-V

```
Loop: lw    t0, 0(s1)      # t0=array element
      add   t0, t0, s2     # add scalar in s2
      sw    t0, 0(s1)     # store result
      addi  s1, s1, -4     # decrement pointer
      bge   s1, zero, Loop # branch s1>=0
```

	ALU/branch	Load/store	cycle
Loop:		lw t0, 0(s1)	1
			2
			3
			4

Note: assuming forwarding is automatically used and there's 1 cycle load-use latency

Scheduling Example

- Schedule this for dual-issue RISC-V

```
Loop: lw    t0, 0(s1)      # t0=array element
      add   t0, t0, s2     # add scalar in s2
      sw    t0, 0(s1)      # store result
      addi  s1, s1, -4     # decrement pointer
      bge   s1, zero, Loop # branch s1>=0
```

	ALU/branch	Load/store	cycle
Loop:		lw t0, 0(s1)	1
			2
	add t0, t0, s2		3
			4

Scheduling Example

- Schedule this for dual-issue RISC-V

```
Loop: lw    t0, 0(s1)      # t0=array element
      add   t0, t0, s2     # add scalar in s2
      sw    t0, 0(s1)     # store result
      addi  s1, s1, -4     # decrement pointer
      bge   s1, zero, Loop # branch s1>=0
```

	ALU/branch	Load/store	cycle
Loop:		lw t0, 0(s1)	1
			2
	add t0, t0, s2		3
		sw t0, 0(s1)	4

Scheduling Example

- Schedule this for dual-issue RISC-V

```
Loop: lw    t0, 0(s1)      # t0=array element
      add   t0, t0, s2     # add scalar in s2
      sw    t0, 0(s1)     # store result
      addi  s1, s1, -4     # decrement pointer
      bge   s1, zero, Loop # branch s1>=0
```

	ALU/branch	Load/store	cycle
Loop:		lw t0, 0(s1)	1
	addi s1, s1, -4		2
	add t0, t0, s2		3
		sw t0, 4(s1)	4

bring addi to cycle 2 to avoid stall, but
need to adjust sw offset

Scheduling Example

- Schedule this for dual-issue RISC-V

```
Loop: lw    t0, 0(s1)      # t0=array element
      add   t0, t0, s2     # add scalar in s2
      sw    t0, 0(s1)      # store result
      addi  s1, s1, -4     # decrement pointer
      bge   s1, zero, Loop # branch s1>=0
```

	ALU/branch	Load/store	cycle
Loop:		lw t0, 0(s1)	1
	addi s1, s1, -4		2
	add t0, t0, s2		3
	bge s1, zero, Loop	sw t0, 4(s1)	4

Scheduling Example

- Schedule this for dual-issue RISC-V

```
Loop: lw    t0, 0(s1)      # t0=array element
      add   t0, t0, s2     # add scalar in s2
      sw    t0, 0(s1)     # store result
      addi  s1, s1, -4     # decrement pointer
      bge   s1, zero, Loop # branch s1>=0
```

	ALU/branch	Load/store	cycle
Loop:	<code>nop</code>	<code>lw t0, 0(s1)</code>	1
	<code>addi s1, s1, -4</code>	<code>nop</code>	2
	<code>add t0, t0, s2</code>	<code>nop</code>	3
	<code>bge s1, zero, Loop</code>	<code>sw t0, 4(s1)</code>	4

Note: `addi` can also be place in cycle 1

$$\text{IPC} = 5/4 = 1.25 \text{ (c.f. peak IPC} = 2\text{)}$$

Loop Unrolling

- “Name dependence” or “anti-dependence”

- Repeated instance of

```
for (i=7; i >= 0;i--)  
a[i]=a[i]+5;
```

```
Loop: lw    t0, 0(s1)  
      add   t0, t0, s2  
      sw    t0, 0(s1)  
      addi  s1, s1, -4  
      bge   s1, zero, Loop
```

```
a[7]=a[7]+5;  
a[6]=a[6]+5;  
a[5]=a[5]+5;  
a[4]=a[4]+5;  
...
```

- The data are independent, no data flow between two sets
- Dependence comes from the reuse of the register name
- We use “loop unrolling” to remove “name dependence”
 - Replicate loop body to expose more parallelism
 - Use different registers per replication (called “register renaming”)
 - Reduces loop-control overhead

Loop Unrolling Example

- Repeat the code in the loop (assuming loop 4 times)

lw t0 , 0(s1) add t0 , t0 , s2 sw t0 , 0(s1)	lw t1 , -4(s1) add t1 , t1 , s2 sw t1 , -4(s1)	lw t2 , -8(s1) add t2 , t2 , s2 sw t2 , -8(s1)	lw t3 , -12(s1) add t3 , t3 , s2 sw t3 , -12(s1)
--	--	--	--

	ALU/branch	Load/store	cycle
Loop:	addi s1 , s1, -16	lw t0 , 0(s1)	1
	nop	lw t1 , 12(s1)	2
	add t0 , t0 , s2	lw t2 , 8(s1)	3
	add t1 , t1 , s2	lw t3 , 4(s1)	4
	add t2 , t2 , s2	sw t0 , 16(s1)	5
	add t3 , t4 , s2	sw t1 , 12(s1)	6
	nop	sw t2 , 8(s1)	7
	bge s1 , zero, Loop	sw t3 , 4(s1)	8

Loop Unrolling Example

- $IPC = 14/8 = 1.75$
 - Closer to 2, but at cost of registers and code size
- How about we loop 3 or 5 iterations, instead of 4?
 - 11/6 if it's 3, 17/10 if it's 5

	ALU/branch	Load/store	cycle
Loop:	addi s1 , s1, -16	lw t0 , 0(s1)	1
	nop	lw t1 , 12(s1)	2
	add t0 , t0 , s2	lw t2 , 8(s1)	3
	add t1 , t1 , s2	lw t3 , 4(s1)	4
	add t2 , t2 , s2	sw t0 , 16(s1)	5
	add t3 , t4 , s2	sw t1 , 12(s1)	6
	nop	sw t2 , 8(s1)	7
	bge s1 , zero, Loop	sw t3 , 4(s1)	8

Dynamic Multiple Issue

- “Superscalar” processors
 - An advanced pipelining techniques that enables the processor to execute more than one instruction per clock cycle **by selecting them during execution.**
- CPU decides whether to issue 0, 1, 2, ... each cycle
 - Avoiding structural and data hazards
- Avoids the need for compiler scheduling
 - Though it may still help
 - Code semantics ensured by the CPU

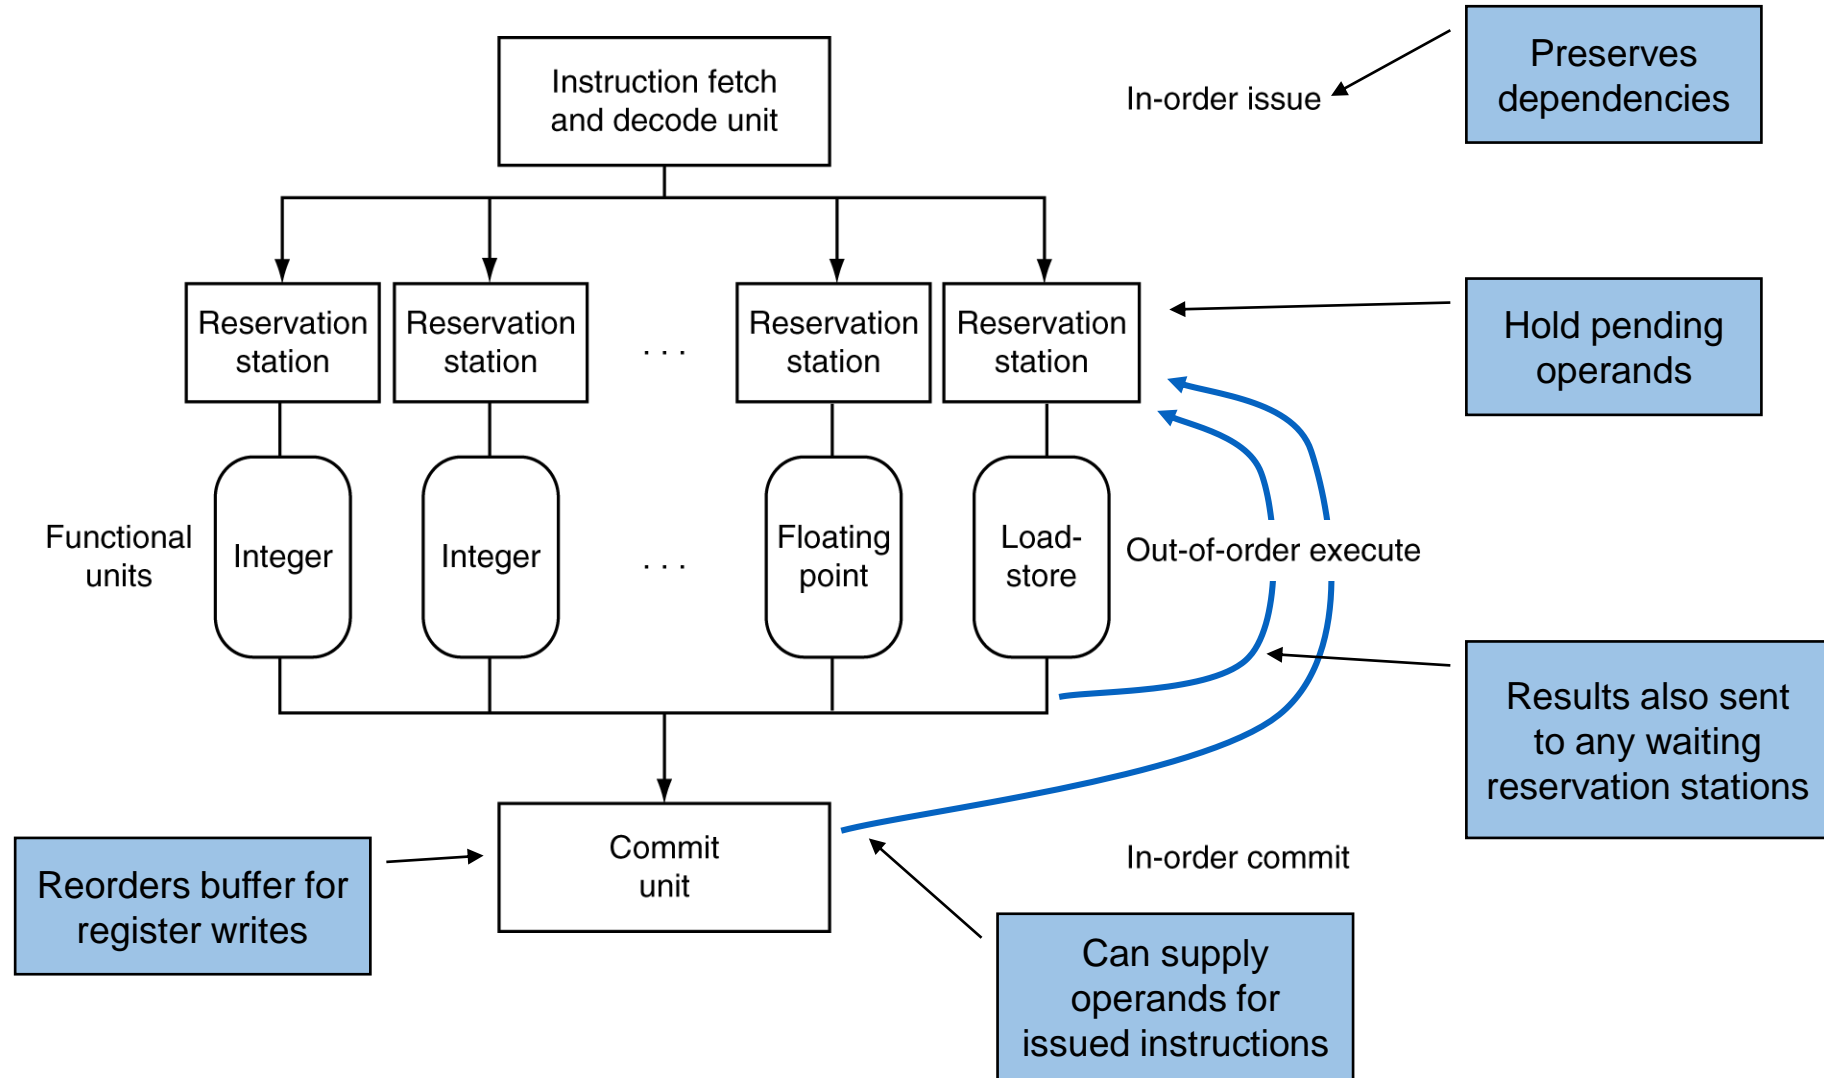
Dynamic Pipeline Scheduling

- Allow the CPU to execute instructions out of order to avoid stalls
 - But commit result to registers in order
- Example

```
lw    x31, 20(x21)
add   x1, x31, x2
sub   x23, x23, x3
andi  x5, x23, 20
```

 - Can start sub while add is waiting for ld

Dynamically Scheduled CPU



Register Renaming

- Reservation stations and reorder buffer effectively provide register renaming
- On instruction issue to reservation station
 - If operand is available in register file or reorder buffer
 - Copied to reservation station
 - No longer required in the register; can be overwritten
 - If operand is not yet available
 - It will be provided to the reservation station by a function unit
 - Register update may not be required

Why Do Dynamic Scheduling?

- Why not just let the compiler schedule code?
- Not all stalls are predictable
 - e.g., cache misses
- Can't always schedule around branches
 - Branch outcome is dynamically determined
- Different implementations of an ISA have different latencies and hazards

Speculation

- “Guess” what to do with an instruction
 - Start operation as soon as possible
 - Check whether guess was right
 - If so, complete the operation
 - If not, roll-back and do the right thing
- Common to static and dynamic multiple issue
- Examples
 - Speculate on branch outcome
 - Roll back if path taken is different
 - Speculate on load
 - Avoid load and cache miss delay
 - Roll back if location is updated

Compiler/Hardware Speculation

- Compiler can reorder instructions
 - e.g., move an instruction across a branch or a load across a store
 - Can include “fix-up” instructions to recover from incorrect guess
- Hardware can look ahead for instructions to execute
 - Buffer results and write to register or memory until it determines they are actually needed
 - Flush buffers on incorrect speculation



Speculation and Exceptions

- What if exception occurs on a speculatively executed instruction?
 - e.g., speculative load before null-pointer check
- Static speculation
 - Can add ISA support for deferring exceptions
- Dynamic speculation
 - Can buffer exceptions until instruction completion (which may not occur)

Does Multiple Issue Work?

- Yes, but not as much as we'd like
- Programs have real dependencies that limit ILP
- Some dependencies are hard to eliminate
 - e.g., pointer aliasing
- Some parallelism is hard to expose
 - Limited window size during instruction issue
- Memory delays and limited bandwidth
 - Hard to keep pipelines full
- Speculation can help if done well

Summary

• Techniques

- Static Multiple issue
- Dynamic Multiple issue
- VLIW
- Superscalar
- Loop unrolling
- Register renaming
- Dynamic scheduling
- Out-of-order execution
- Speculation

	Static multiple issue	Dynamic multiple issue
Decision made by	Compiler (software)	Processor (hardware)
Also called	Very long instruction word (VLIW)	Superscalar
Ways to remove hazard	Loop unrolling Register renaming Speculation	Dynamic scheduling Out-of-order execution Register renaming Speculation

Power Efficiency


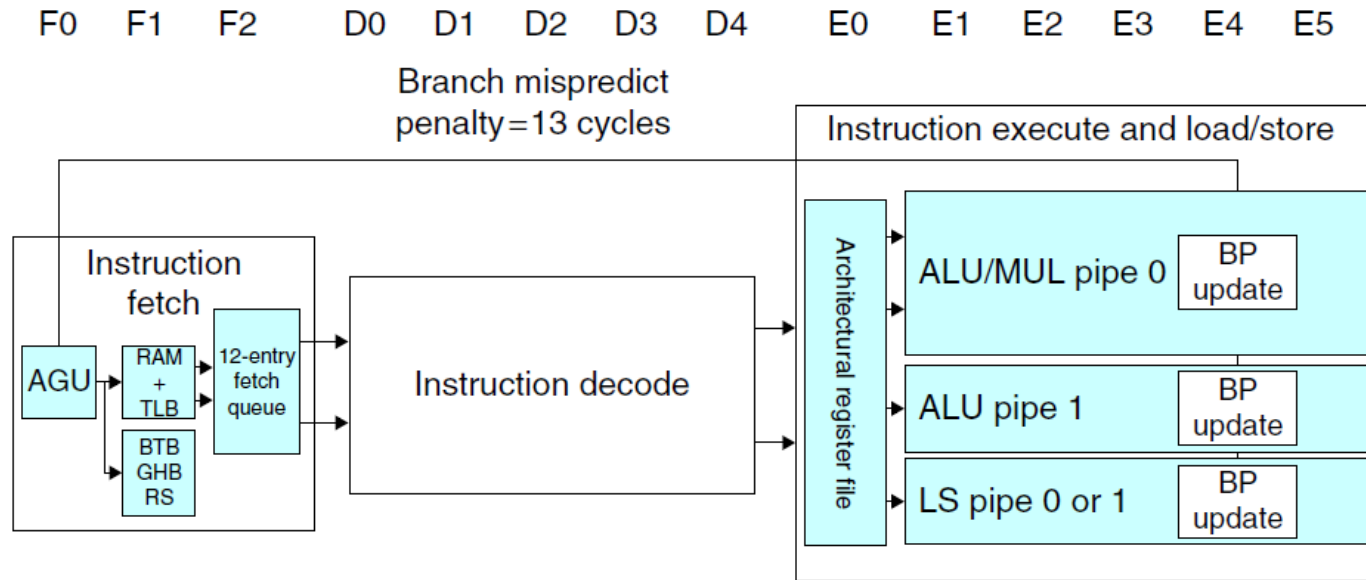
- Complexity of dynamic scheduling and speculations requires power
- Multiple simpler cores may be better

Microprocessor	Year	Clock Rate	Pipeline Stages	Issue width	Out-of-order/ Speculation	Cores	Power
i486	1989	25MHz	5	1	No	1	5W
Pentium	1993	66MHz	5	2	No	1	10W
Pentium Pro	1997	200MHz	10	3	Yes	1	29W
P4 Willamette	2001	2000MHz	22	3	Yes	1	75W
P4 Prescott	2004	3600MHz	31	3	Yes	1	103W
Core	2006	2930MHz	14	4	Yes	2	75W
UltraSparc III	2003	1950MHz	14	4	No	1	90W
UltraSparc T1	2005	1200MHz	6	1	No	8	70W


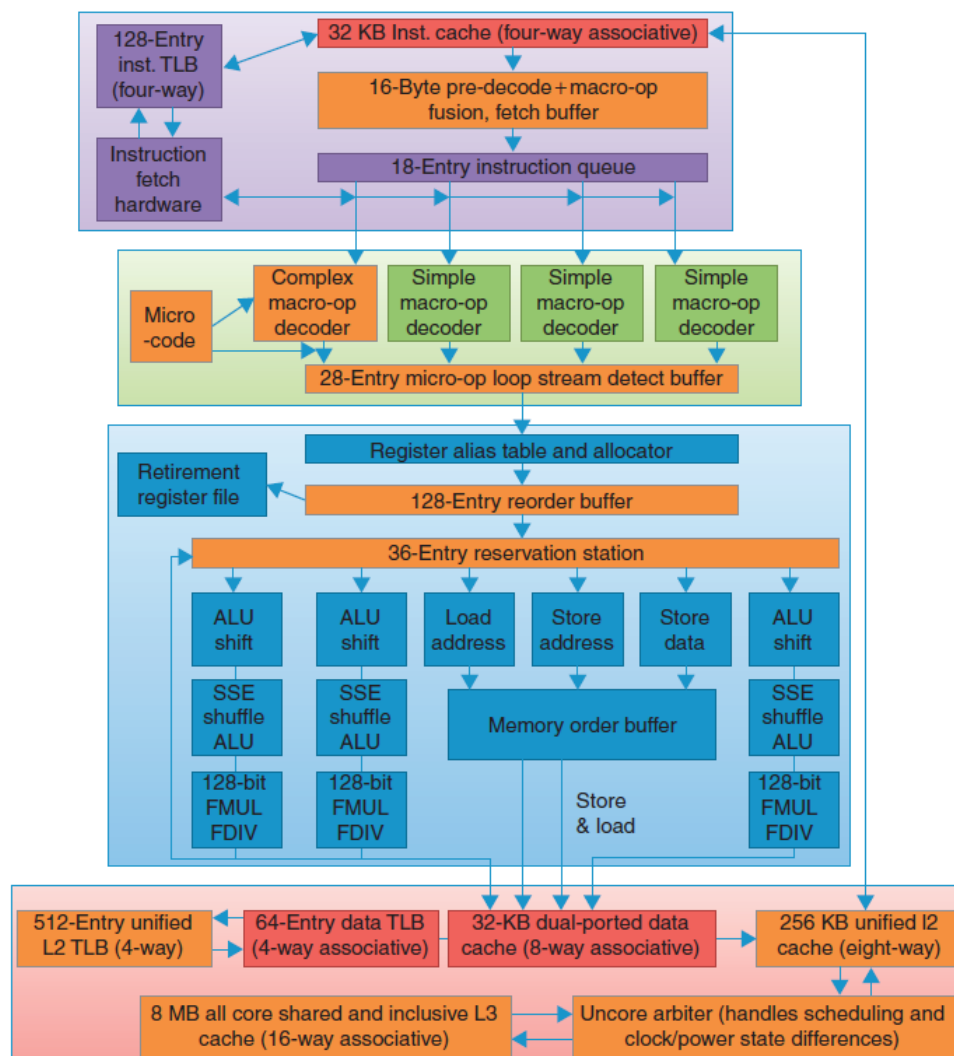
Cortex A8 and Intel i7

Processor	ARM A8	Intel Core i7 920
Market	Personal Mobile Device	Server, cloud
Thermal design power	2 Watts	130 Watts
Clock rate	1 GHz	2.66 GHz
Cores/Chip	1	4
Floating point?	No	Yes
Multiple issue?	Dynamic	Dynamic
Peak instructions/clock cycle	2	4
Pipeline stages	14	14
Pipeline schedule	Static in-order	Dynamic out-of-order with speculation
Branch prediction	2-level	2-level
1 st level caches/core	32 KiB I, 32 KiB D	32 KiB I, 32 KiB D
2 nd level caches/core	128-1024 KiB	256 KiB
3 rd level caches (shared)	-	2- 8 MB

ARM Cortex-A8 Pipeline & Performance



Core i7 Pipeline & Performance



Fallacies

- Pipelining is easy (!)
 - The basic idea is easy
 - The devil is in the details
 - e.g., detecting data hazards
- Pipelining is independent of technology
 - So why haven't we always done pipelining?
 - More transistors make more advanced techniques feasible
 - Pipeline-related ISA design needs to take account of technology trends

Pitfalls

- Poor ISA design can make pipelining harder
 - e.g., complex instruction sets (VAX, IA-32)
 - Significant overhead to make pipelining work
 - IA-32 micro-op approach
 - e.g., complex addressing modes
 - Register update side effects, memory indirection
 - e.g., delayed branches
 - Advanced pipelines have long delay slots

Concluding Remarks

- ISA influences design of datapath and control
- Datapath and control influence design of ISA
- Pipelining improves instruction throughput using parallelism
 - More instructions completed per second
 - Latency for each instruction not reduced
- Hazards: structural, data, control
- Multiple issue and dynamic scheduling (ILP)
 - Dependencies limit achievable parallelism
 - Complexity leads to the power wall